




MI2022 HW2

授課教師: 鄭文皇教授

助教: 謝嘉晉、翁紹恩

- 60k images in total with 100 class (balanced class distribution)
- Each class: 500 training images, 100 testing images.
- Images size: 32x32x3 (RGB)



Problem description – CIFAR100

<https://www.cs.toronto.edu/~kriz/cifar.html>

Train a classification model with limited size, 100MB. Then test your model on testing data with report showing your method to prove the provided code (baseline).

Data: cifar100 from Torchvision.

Training data:

```
trainset = torchvision.datasets.CIFAR100(root='./data', train=True,  
                                         download=True, transform=transform)
```

Testing data with accuracy:

```
testset = torchvision.datasets.CIFAR100(root='./data', train=False,  
                                         download=True, transform=transform)
```

Problem description

Test your model with provided test code segment:

```
correct = 0
total = 0
# since we're not training, we don't need to calculate the gradients for our outputs
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images = images.to(device)
        labels = labels.to(device)
        # calculate outputs by running images through the network
        outputs = model(images)
        # the class with the highest energy is what we choose as prediction
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f'Accuracy of the network on the 10000 test images: {100 * correct // total} %')
```

Then save your model to upload E3:

`torch.save(model, “./saved/model.pth”)`

Upload and Test

30% test ranking

Main Component:

1. Dataset preprocessing: data augmentation
2. Model architecture: other efficient model (less than 100 MB for entire saved model)
3. Loss function: focal loss, consistency loss.
4. Optimizer: learning rate scheduler
5. Training strategy (extra data): self / semi / transfer learning, domain adaptation
6. Other

Improve the code and **report** your techniques

70% report

1. Dataset preprocessing: data augmentation

```
import torch
import torchvision
import torchvision.transforms as tra
nsforms
import torchvision.models as models
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

device = torch.device('cuda' if torc
h.cuda.is_available() else 'cpu')
print(device)
```

Give me your mean & std in report
for testing if changing them!

```
mean = (0.5071, 0.4867, 0.4408)
std = (0.2675, 0.2565, 0.2761)

train_transform = transforms.Compose(
    [transforms.RandomHorizontalFlip(p=0.5),
     transforms.ToTensor(),
     transforms.Normalize(mean, std)]) # calculte yourself

test_transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize(mean, std)]) # calculte yourself
```

Error: CUDA out of memory!

```
batch_size = 32
num_classes = 100 # check

trainset = torchvision.datasets.CIFAR100(root='./data', train=True,
                                         download=True, transform=train_transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR100(root='./data', train=False,
                                         download=True, transform=test_transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=2)
```

Supporting Code

<https://pytorch.org/vision/stable/transforms.html>

2. Model architecture: other efficient model (1/2)

Define by your self or clone other code

```
class Toy_CNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, num_classes)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

off-the-shelf model

ResNet

`resnet18` ([pretrained, progress])

`resnet34` ([pretrained, progress])

`resnet50` ([pretrained, progress])

```
model = models.resnet50(pretrained=True)
model.fc = torch.nn.Linear(2048, num_classes)
```


**ResNet50 is small than 100MB*

Supporting Code

PyTorch model:
<https://pytorch.org/vision/stable/models.html>

2. Model architecture: other efficient model (2/2)

```
model = models.mobilenet_v3_small(pretrained=True)
model.classifier[3] = torch.nn.Linear(1024, num_classes)
```



```
print(model)
```

⋮

```
(avgpool): AdaptiveAvgPool2d(output_size=1)
(classifier): Sequential(
  (0): Linear(in_features=576, out_features=1024, bias=True)
  (1): Hardswish()
  (2): Dropout(p=0.2, inplace=True)
  (3): Linear(in_features=1024, out_features=100, bias=True)
)
```



```
model = models.mobilenet_v3_small(pretrained=True)
model.classifier[3] = torch.nn.Linear(1024, num_classes)
```

Change the last layer by code

Supporting Code

`torch.nn.Linear:`
<https://pytorch.org/docs/stable/generated/torch.nn.Linear.html>

3. Loss function

4. Optimizer

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.Adam(model.parameters(), lr=0.0001, weight_decay=1e-4)
```

more PyTorch built-in loss function:

<https://pytorch.org/docs/stable/nn.html#loss-functions>

more PyTorch built-in optimizer:

<https://pytorch.org/docs/stable/optim.html>

Supporting Code

5. Training strategy (1/2)

Don't use
`torch.save(model.state_dict(),
save_path)`

```
total_epoch = 20
print_per_iteration = 100
save_path = './model.pth'
major_class = 10

for epoch in range(total_epoch): # loop over the dataset multiple times
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)
        # zero the parameter gradients
        optimizer.zero_grad()
        # forward + backward + optimize
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        if (i+1) % print_per_iteration == 0: # print every 2000 mini-batches
            print(f'[ep {epoch + 1}][{i + 1:5d}/{len(trainloader):5d}] loss: {loss.item():.3f}')
    torch.save(model, save_path)
```

Supporting Code

5. Training strategy (2/2)

(extra data): self / semi / transfer learning, long tail distributions technique

- FixMatch(Semi-S): <https://github.com/kekmodel/FixMatch-pytorch>
- BYOR(Self-S): <https://github.com/lucidrains/byol-pytorch>
- Pseudo long-tail techniques:

```
criterion = nn.CrossEntropyLoss(reduction='none')
optimizer = optim.Adam(model.parameters(), lr=0.0001,
                        weight_decay=1e-4)
def get_major_class(outputs):
    pass
```

Just a sample !

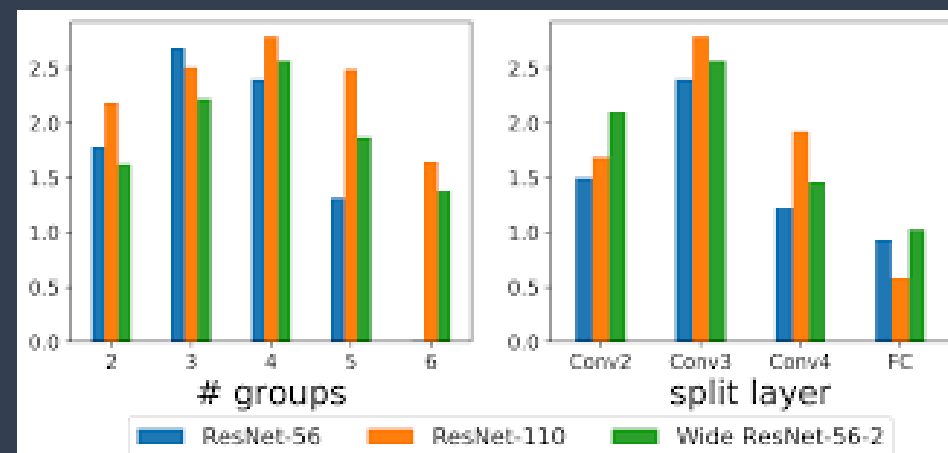

Requires more modification to work !

```
for epoch in range(total_epoch): # loop over the dataset multiple times
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)
        # zero the parameter gradients
        optimizer.zero_grad()
        # forward + backward + optimize
        outputs = model(inputs)
        with torch.no_grad():
            major_class = get_major_class(outputs) # no_grad
            mask = (torch.max(outputs, dim=1)[1] == major_class).float()
        loss = (criterion(outputs, labels) * mask).mean()
        loss.backward()
        optimizer.step()
```

Supporting Code

6. Other

調參數大師
(lr, batch_size ...)



Supporting Code

The screenshot displays the Google Colab web interface. The top menu bar includes 'File', 'Edit', 'View', 'Insert', 'Runtime', 'Tools', and 'Help'. The 'Runtime' menu is open, showing options like 'Run all', 'Run before', 'Run the focused cell', 'Run selection', 'Run after', 'Interrupt execution', 'Restart runtime', 'Restart and run all', 'Factory reset runtime', 'Change runtime type' (highlighted with a red circle), 'Manage sessions', and 'View runtime logs'. A blue arrow points from the 'Change runtime type' option to the 'Notebook settings' dialog box on the right.

The 'Notebook settings' dialog box is open, showing the 'Hardware accelerator' set to 'GPU' (indicated by a red checkmark). Below this, there is a message: 'To get the most out of Colab, avoid using a GPU unless you need one. [Learn more](#)'. There are also checkboxes for 'Background execution' and 'Omit code cell output when saving this notebook'. At the bottom of the dialog are 'Cancel' and 'Save' buttons.

In the background, the Colab interface shows a file explorer on the left with folders 'data' and 'sample_data', and files 'model.pth' and 'model_state.pth'. The main editor area contains Python code for data augmentation using PyTorch's torchvision library.

Colab for quick starting

https://colab.research.google.com/?utm_source=scs-index

Upload file

1. Trained model: hw2_ID.pth

- check the test accuracy by the code ☞
- Less than 100MB (otherwise will fail to upload the model.pth to E3)

2. Report: hw2_ID.pdf

3. (optional) hw2_ID.txt

```
mean = (0.5, 0.5, 0.5)
std = (0.5, 0.5, 0.5)

train_transform = transforms.Compose(
    [transforms.RandomHorizontalFlip(p=0.5),
     transforms.ToTensor(),
     transforms.Normalize(mean, std)]) # calculate yourself

test_transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize(mean, std)]) # calculate yourself

batch_size = 32
num_classes = 100 # check

trainset = torchvision.datasets.CIFAR100(root='./data', train=True,
                                         download=True, transform=train_transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR100(root='./data', train=False,
                                         download=True, transform=test_transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=2)
```

hw2_309510151.txt - 記事本

檔案(F) 編輯(E) 格式(O) 檢視(V) 說明

0.5 0.5 0.5
0.24 0.24 0.24

Deadline 4/10 23:55